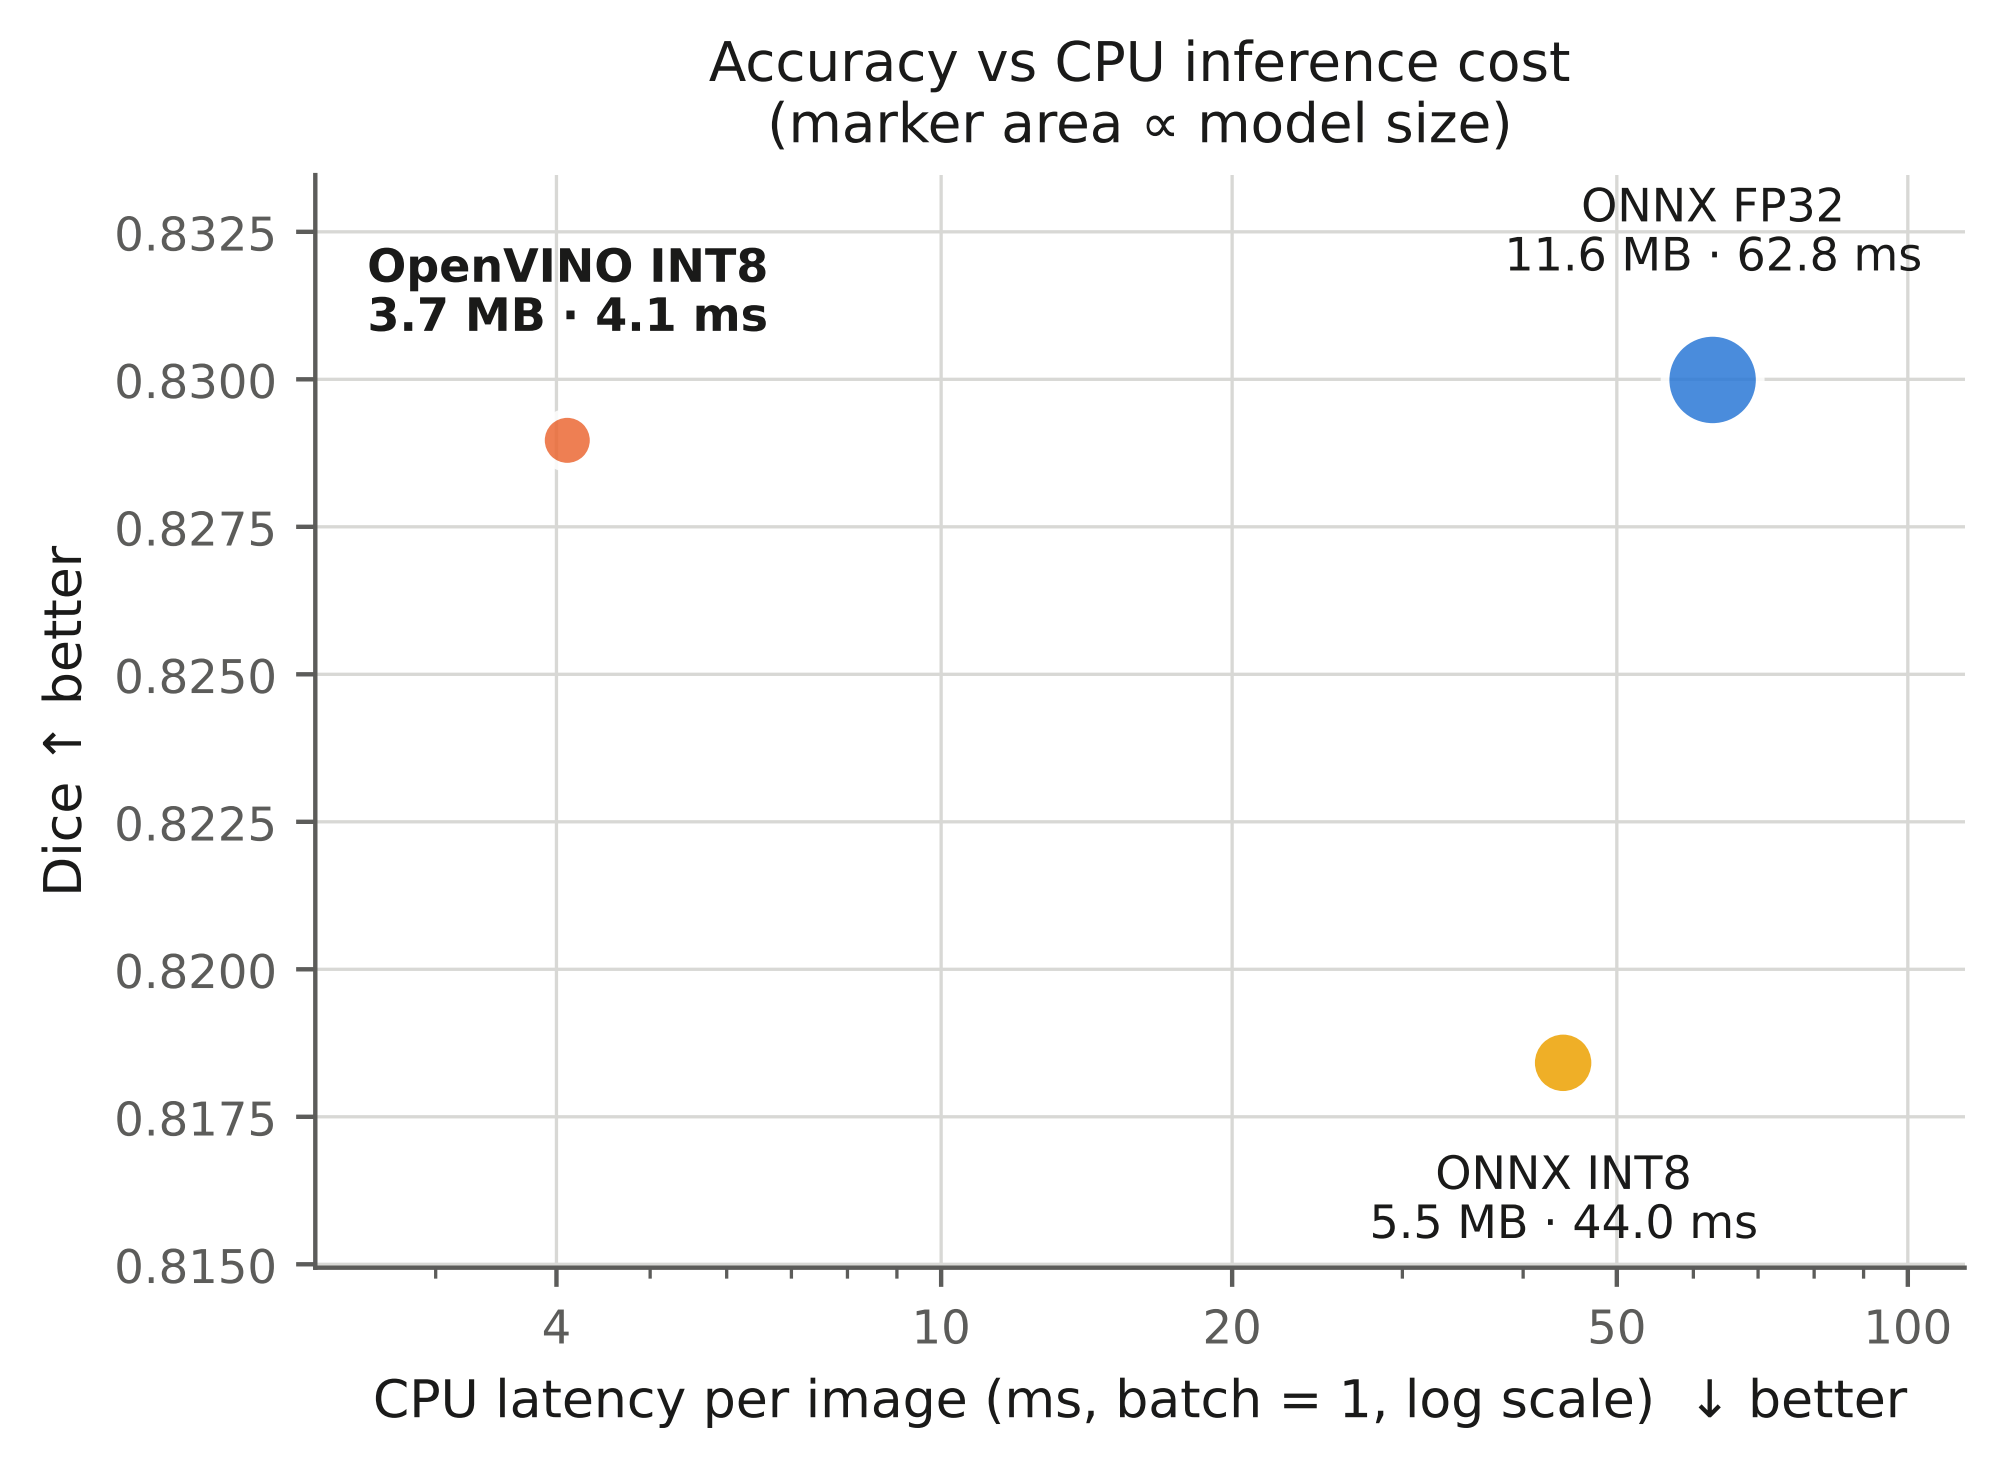
Accuracy vs CPU inference cost
(marker area ∝ model size)
OpenVINO INT8
3.7 MB · 4.1 ms
ONNX FP32
11.6 MB · 62.8 ms
ONNX INT8
5.5 MB · 44.0 ms
Dice ↑ better
0.8325
0.8300
0.8275
0.8250
0.8225
0.8200
0.8175
0.8150
4
10
20
50
100
CPU latency per image (ms, batch = 1, log scale) ↓ better