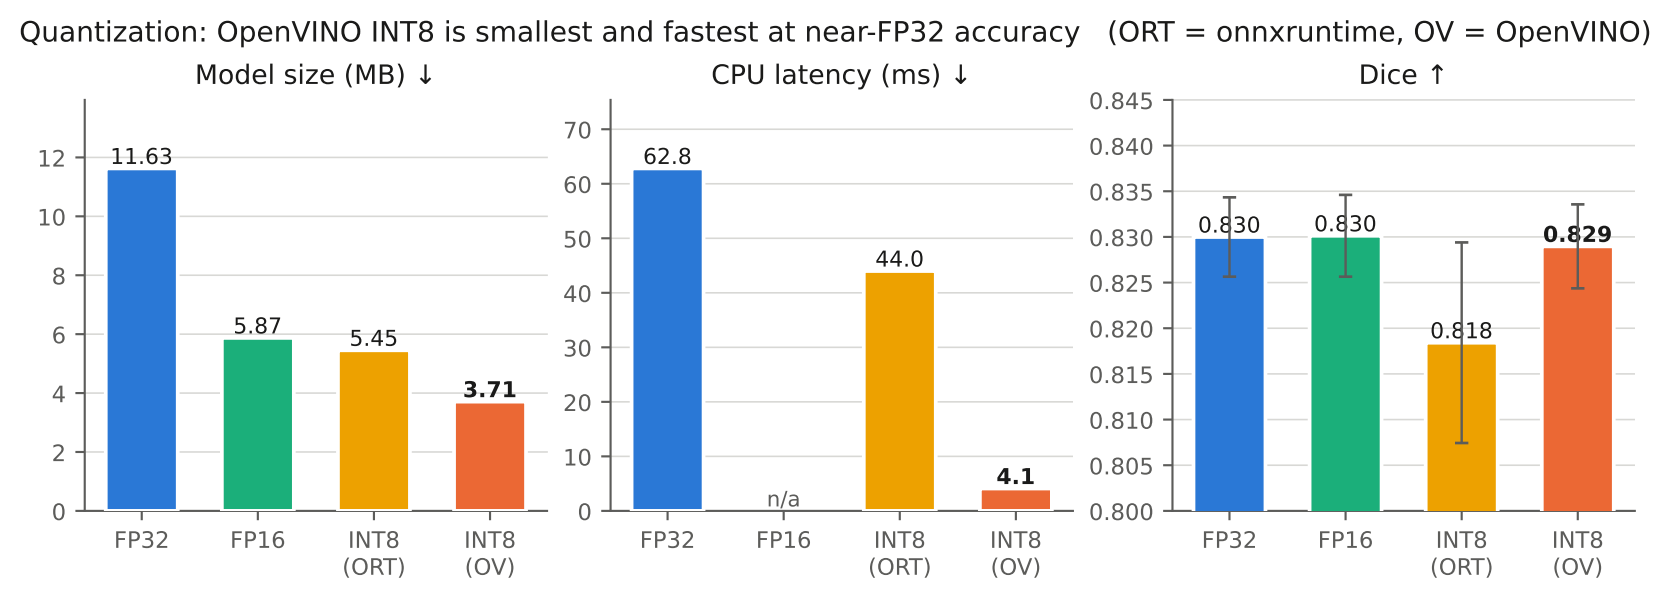
Quantization: OpenVINO INT8 is smallest and fastest at near-FP32 accuracy (ORT = onnxruntime, OV = OpenVINO)
Model size (MB) ↓
11.63
5.87
5.45
3.71
FP32
FP16
INT8 (ORT)
INT8 (OV)
CPU latency (ms) ↓
62.8
n/a
44.0
4.1
Dice ↑
0.830
0.830
0.818
0.829